

# 6.036 Machine learning

Review lecture

# Final Exam (1:30pm, May 18)

- 3h exam, i.e., more comprehensive than midterm
- closed book, you cannot bring a cheatsheet
- we will give you a cheatsheet at the exam (and post it online ahead of time, no later than Monday)
- it covers the whole course material but has roughly speaking  $\frac{2}{3}$  emphasis on the material since midterm
- questions are of the same kind as in the midterm
- there can be broader questions from guest lectures
- previous finals (as practice problems) already available on Stellar

# Review sessions

- Two "office hour style" (come ask questions) sessions

7-9pm Monday May 16 in 34-101

7-9pm Tuesday May 17 in 34-101

# Material before the midterm

- Linear classification, separation
- On-line linear classifiers
- Maximum margin separation; SVM, Pegasos
- Linear regression
- Collaborative filtering
- Non-linear classifiers, kernels
- Neural networks, deep learning, back-propagation
- Recurrent neural networks

# Material since the midterm

- Generalization, VC-dimension
  - Clustering (k-means)
  - Mixture models
  - EM algorithm
  - Bayesian Networks
  - Hidden Markov Models
  - MDPs, Reinforcement Learning
- + guest lectures

# Generalization, VC-dim

- Need to understand (qualitatively) how the difference between training and test errors depend on the number of training examples ( $n$ ) and the (log of the) number of classifiers considered ( $\log|H|$ )
- Shattering and VC dimension, calculating the VC dimension of a simple set of classifiers
- Note that there are many ways to understand generalization: how log-likelihoods of training and test data for a mixture model behave as we change the number of mixture components

# Clustering

- K-means clustering
- How the algorithm behaves
  - with a different initialization
  - with a different choice of metric

# Mixture models

- Definition of a mixture of spherical Gaussians, what the parameters are, what they control
- How to sample from a mixture model
- Types of data that the mixture model is designed to model well (e.g., overlapping clusters)
- We use the log-likelihood of the data to measure how well a mixture model agrees with the data. Need to understand (qualitatively) how this log-likelihood behaves if we change the mixture parameters



# EM algorithm for mixtures

- E-step of the EM algorithm (cf. assignment step in the k-means algorithm)
- M-step of the EM algorithm (cf. update of cluster centroids in k-means)
- Should be able to assess by hand how the algorithm modifies the mixture in simple cases (cf. previous exam questions)
- Need to understand the effect of different initializations on the resulting model
  - e.g., if all the mixture components are initially the same
  - e.g., same means and mixing proportions, different variances

# Bayesian networks

- Understand "arcs" as dependencies, interpret multiple incoming "arcs"
- The notion of probability tables (e.g., one variable depending on two others)
- How to write the distribution over all the variables as a product of simpler terms based on the graph
- Basic independence statements
- Be able to draw Mixture models, Markov, and Hidden Markov Models as Bayesian networks

# Hidden Markov Models

- Need to understand them "as models" (cf. previous exam questions)
- If given the initial state distribution, transition probability matrix, and the emission (output) distribution, you should be able to evaluate
  - the probability of a particular sequence of states and observations
  - the probability of a sequence of observations (not knowing the states)
  - the most likely sequence of hidden states corresponding to a sequence of observations
  - the probability that the state at time  $t$  took a particular value if given the observations

# Markov Decision Problems (MDP)

- Definition in terms of states, actions, rewards, transition probabilities, discount factor
- What value function is, how to calculate it through value iteration (or from Q-values)
- Q-values (what they mean), how to get them from Q-value iteration (or based on the value function)
- policy derived from Q-values; optimal policy
- For example: should be able to calculate how the policy derived from Q-values changes as a function of Q-value iterations (cf. previous exam questions)